

# Refactored Capstone Notebook

## What Changed

The core logic, features, and models remain the same, but the following improvements were made:

### Key Improvements

- All preprocessing moved into **ColumnTransformer + Pipeline**
- **SMOTE** applied *inside* pipelines (no data leakage)
- Consistent **Stratified Cross-Validation**
- Unified evaluation across models
- Entire pipelines saved for deployment

## 1. Imports

```
In [1]: pip install optuna
```

```
Collecting optuna
  Downloading optuna-4.7.0-py3-none-any.whl.metadata (17 kB)
Requirement already satisfied: alembic>=1.5.0 in /usr/local/lib/python3.12/dist-packages (from optuna) (1.18.3)
Collecting colorlog (from optuna)
  Downloading colorlog-6.10.1-py3-none-any.whl.metadata (11 kB)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (from optuna) (2.0.2)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from optuna) (26.0)
Requirement already satisfied: sqlalchemy>=1.4.2 in /usr/local/lib/python3.12/dist-packages (from optuna) (2.0.46)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-packages (from optuna) (4.67.3)
Requirement already satisfied: PyYAML in /usr/local/lib/python3.12/dist-packages (from optuna) (6.0.3)
Requirement already satisfied: Mako in /usr/local/lib/python3.12/dist-packages (from alembic>=1.5.0->optuna) (1.3.10)
Requirement already satisfied: typing-extensions>=4.12 in /usr/local/lib/python3.12/dist-packages (from alembic>=1.5.0->optuna) (4.15.0)
Requirement already satisfied: greenlet>=1 in /usr/local/lib/python3.12/dist-packages (from sqlalchemy>=1.4.2->optuna) (3.3.1)
Requirement already satisfied: MarkupSafe>=0.9.2 in /usr/local/lib/python3.12/dist-packages (from Mako->alembic>=1.5.0->optuna) (3.0.3)
  Downloading optuna-4.7.0-py3-none-any.whl (413 kB)
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 413.9/413.9 kB 10.7 MB/s eta 0:00:00
  Downloading colorlog-6.10.1-py3-none-any.whl (11 kB)
  Installing collected packages: colorlog, optuna
  Successfully installed colorlog-6.10.1 optuna-4.7.0
```

```
In [2]: import pandas as pd
import numpy as np
```

```

import matplotlib.pyplot as plt
import seaborn as sns
import optuna
import xgboost as xgb
import lightgbm as lgb
import joblib
import warnings
warnings.filterwarnings("ignore")

from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder, StandardScaler, F
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_sco
from sklearn.feature_selection import mutual_info_classif, SelectKBest
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier,
from sklearn.svm import SVC
from scipy.stats import f_oneway, ttest_ind
from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.impute import SimpleImputer
from imblearn.over_sampling import SMOTE

```

## 2. Data Loading

In [3]: # Mount Google Drive to access files stored in my google Drive

```

from google.colab import drive
drive.mount('/content/drive')

```

Mounted at /content/drive

In [4]: def load\_data(train\_path, test\_path):
 """Load train and test datasets from given file paths."""
 train\_df = pd.read\_csv(train\_path)
 test\_df = pd.read\_csv(test\_path)
 return train\_df, test\_df

df\_train, df\_test = load\_data('/content/drive/MyDrive/Capstone Project Data/trai
 '/content/drive/MyDrive/Capstone Project Data/test

# Display samples
print("Train Sample:")
display(df\_train.head())
print("Test Sample:")
display(df\_test.head())

# Display shapes
print(f"Train shape: {df\_train.shape}")
print(f"Test shape: {df\_test.shape}")

Train Sample:

	<b>id</b>	<b>day</b>	<b>pressure</b>	<b>maxtemp</b>	<b>temparature</b>	<b>mintemp</b>	<b>dewpoint</b>	<b>humidity</b>	<b>cloud</b>	<b>sun</b>
<b>0</b>	0	1	1017.4	21.2	20.6	19.9	19.4	87.0	88.0	
<b>1</b>	1	2	1019.5	16.2	16.9	15.8	15.4	95.0	91.0	
<b>2</b>	2	3	1024.1	19.4	16.1	14.6	9.3	75.0	47.0	
<b>3</b>	3	4	1013.4	18.1	17.8	16.9	16.8	95.0	95.0	
<b>4</b>	4	5	1021.8	21.3	18.4	15.2	9.6	52.0	45.0	

◀ ▶ Test Sample:

	<b>id</b>	<b>day</b>	<b>pressure</b>	<b>maxtemp</b>	<b>temparature</b>	<b>mintemp</b>	<b>dewpoint</b>	<b>humidity</b>	<b>cloud</b>	<b>s</b>
<b>0</b>	2190	1	1019.5	17.5	15.8	12.7	14.9	96.0	99.0	
<b>1</b>	2191	2	1016.5	17.5	16.5	15.8	15.1	97.0	99.0	
<b>2</b>	2192	3	1023.9	11.2	10.4	9.4	8.9	86.0	96.0	
<b>3</b>	2193	4	1022.9	20.6	17.3	15.2	9.5	75.0	45.0	
<b>4</b>	2194	5	1022.2	16.1	13.8	6.4	4.3	68.0	49.0	

◀ ▶ Train shape: (2190, 13)

Test shape: (730, 12)

### 3. Basic Info & Summary

```
In [5]: # Check data types
print(df_train.info())

# Summary statistics
print(df_train.describe())

# Check for missing values
print(df_train.isnull().sum())
print(df_test.isnull().sum())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2190 entries, 0 to 2189
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               2190 non-null    int64  
 1   day              2190 non-null    int64  
 2   pressure         2190 non-null    float64 
 3   maxtemp          2190 non-null    float64 
 4   temparature      2190 non-null    float64 
 5   mintemp          2190 non-null    float64 
 6   dewpoint         2190 non-null    float64 
 7   humidity          2190 non-null    float64 
 8   cloud             2190 non-null    float64 
 9   sunshine          2190 non-null    float64 
 10  windddirection   2190 non-null    float64 
 11  windspeed         2190 non-null    float64 
 12  rainfall          2190 non-null    int64  
dtypes: float64(10), int64(3)
memory usage: 222.6 KB
None
      id        day      pressure     maxtemp  temparature \
count  2190.000000  2190.000000  2190.000000  2190.000000  2190.000000
mean   1094.500000  179.948402  1013.602146  26.365799   23.953059
std    632.342866  105.203592   5.655366   5.654330   5.222410
min    0.000000    1.000000   999.000000  10.400000   7.400000
25%   547.250000   89.000000  1008.600000  21.300000  19.300000
50%   1094.500000  178.500000  1013.000000  27.800000  25.500000
75%   1641.750000  270.000000  1017.775000  31.200000  28.400000
max   2189.000000  365.000000  1034.600000  36.000000  31.500000

      mintemp      dewpoint     humidity      cloud      sunshine \
count  2190.000000  2190.000000  2190.000000  2190.000000  2190.000000
mean   22.170091   20.454566   82.036530   75.721918   3.744429
std    5.059120   5.288406   7.800654   18.026498   3.626327
min    4.000000   -0.300000   39.000000   2.000000   0.000000
25%   17.700000   16.800000   77.000000   69.000000   0.400000
50%   23.850000   22.150000   82.000000   83.000000   2.400000
75%   26.400000   25.000000   88.000000   88.000000   6.800000
max   29.800000   26.700000   98.000000  100.000000  12.100000

      windddirection      windspeed      rainfall
count  2190.000000  2190.000000  2190.000000
mean   104.863151   21.804703   0.753425
std    80.002416   9.898659   0.431116
min    10.000000   4.400000   0.000000
25%   40.000000   14.125000   1.000000
50%   70.000000   20.500000   1.000000
75%   200.000000  27.900000   1.000000
max   300.000000  59.500000   1.000000


      id      day      pressure     maxtemp  temparature \
id      0
day      0
pressure 0
maxtemp 0
temparature 0
mintemp 0
dewpoint 0
humidity 0
cloud 0
sunshine 0

```

```
winddirection    0  
windspeed       0  
rainfall        0  
dtype: int64  
id              0  
day             0  
pressure        0  
maxtemp         0  
temparature     0  
mintemp         0  
dewpoint        0  
humidity        0  
cloud            0  
sunshine         0  
winddirection   1  
windspeed        0  
dtype: int64
```

## 4.Target Variable Analysis

```
In [6]: # Visualize the distribution of the target variable 'rainfall'  
  
sns.countplot(x='rainfall', data=df_train)  
plt.title("Rainfall Distribution")  
plt.show()  
  
# Percentage of each class  
print(df_train['rainfall'].value_counts(normalize=True))
```




```
rainfall
1    0.753425
0    0.246575
Name: proportion, dtype: float64
```

In [7]: # Identify all numeric features for analysis

```
num_cols = df_train.select_dtypes(include='number').columns.tolist()
num_cols.remove('rainfall') # exclude target
```

In [8]: # Histograms

```
df_train[num_cols].hist(figsize=(12, 10), bins=20)
plt.tight_layout()
plt.show()
```



In [9]: # Check skewness of numeric features

# Helps decide if transformations are needed (for Linear/logistic models)  
# Strongly skewed features: cloud, dewpoint (consider transforming if necessary)

```
df_train[num_cols].skew()
```


Out[9]:



	0
<b>id</b>	0.000000
<b>day</b>	0.030615
<b>pressure</b>	0.284062
<b>maxtemp</b>	-0.490890
<b>temparature</b>	-0.557471
<b>mintemp</b>	-0.649179
<b>dewpoint</b>	-0.997889
<b>humidity</b>	-0.561541
<b>cloud</b>	-1.339274
<b>sunshine</b>	0.639871
<b>winddirection</b>	0.708063
<b>windspeed</b>	0.769390

**dtype:** float64


In [10]:

```
# Boxplots for outliers
for col in num_cols:
    sns.boxplot(x=df_train[col])
    plt.title(col)
    plt.show()
```







maxtemp




temparature




mintemp




dewpoint




humidity




cloud




sunshine



winddirection



windspeed



```
In [11]: outlier_summary = []


for col in num_cols:
    Q1 = df_train[col].quantile(0.25)
    Q3 = df_train[col].quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
    outliers = df_train[(df_train[col] < lower) | (df_train[col] > upper)][col]
    outlier_summary.append({
        'Feature': col,
        'Num_Outliers': len(outliers),
        'Min_Outlier': outliers.min() if len(outliers) > 0 else None,
        'Max_Outlier': outliers.max() if len(outliers) > 0 else None
    })

outlier_df = pd.DataFrame(outlier_summary).sort_values(by='Num_Outliers', ascending=False)
display(outlier_df)
```

	Feature	Num_Outliers	Min_Outlier	Max_Outlier
8	cloud	129	2.0	40.0
7	humidity	28	39.0	60.0
11	windspeed	28	50.6	59.5
6	dewpoint	26	-0.3	4.4
2	pressure	4	1032.3	1034.6
5	mintemp	1	4.0	4.0
0	id	0	NaN	NaN
1	day	0	NaN	NaN
3	maxtemp	0	NaN	NaN
4	temparature	0	NaN	NaN
9	sunshine	0	NaN	NaN
10	winddirection	0	NaN	NaN

In [12]: *# Visualizing columns with the most outliers*

```
plt.figure(figsize=(10,5))
sns.barplot(data=outlier_df, x='Feature', y='Num_Outliers', palette='viridis')
plt.title("Number of Outliers per Numeric Feature")
plt.xticks(rotation=45)
plt.ylabel("Count of Outliers")
plt.show()
```




In [13]: *# Correlation heatmap*

```
plt.figure(figsize=(12,8))
```

```

corr = df_train[num_cols + ['rainfall']].corr()
mask = np.triu(np.ones_like(corr, dtype=bool))
sns.heatmap(corr, annot=True, cmap='coolwarm', mask=mask)
plt.title("Feature Correlation Heatmap")
plt.show()

```




In [14]: # Highlighting Strong correlations

```

strong_corr = corr.abs() > 0.8
sns.heatmap(corr, annot=True, cmap='coolwarm', mask=~strong_corr)
plt.title("Strong Feature Correlations (|r|>0.8)")
plt.show()

```







```
In [15]: cat_cols = df_train.select_dtypes(include='object').columns.tolist()

for col in cat_cols:
    sns.countplot(y=col, data=df_train, order=df_train[col].value_counts().index)
    plt.title(f"{col} Distribution")
    plt.show()
```


## 5. Feature Relationships with Target

```
In [16]: # Numerical vs Target
for col in num_cols:
    sns.boxplot(x='rainfall', y=col, data=df_train)
    plt.title(f"{col} vs Rainfall")
    plt.show()
```










temparature vs Rainfall




mintemp vs Rainfall




**dewpoint vs Rainfall****humidity vs Rainfall**



## winddirection vs Rainfall



## windspeed vs Rainfall



```
In [17]: # Categorical vs Target
for col in cat_cols:
    sns.countplot(x=col, hue='rainfall', data=df_train)
    plt.title(f"{col} vs Rainfall")
    plt.show()
```


```
In [18]: def day_to_month(day):
    if day <= 31: return 1
    elif day <= 59: return 2
    elif day <= 90: return 3
    elif day <= 120: return 4
    elif day <= 151: return 5
    elif day <= 181: return 6
    elif day <= 212: return 7
    elif day <= 243: return 8
    elif day <= 273: return 9
    elif day <= 304: return 10
    elif day <= 334: return 11
    else: return 12

df_train["month"] = df_train["day"].apply(day_to_month)
```

```
In [19]: monthly_rainfall = df_train.groupby("month")["rainfall"].mean()
```

```
In [20]: import calendar
monthly_rainfall.index = monthly_rainfall.index.map(lambda x: calendar.month_abbr[x])
```

```
In [21]: plt.figure(figsize=(10,5))
plt.plot(monthly_rainfall.index, monthly_rainfall.values)
plt.grid(True, linestyle='--', alpha=0.5)
plt.xlabel("Month")
plt.ylabel("Rainfall Rate")
plt.title("Average Rainfall by Month")
plt.show()
```



**Outcome: After this EDA, you'll know:**

Which features are useful or need transformation

If target is imbalanced

Outliers or missing values to handle

Relationships that could inspire new features

## 6.Train-Test Split

```
In [22]: # Dropping 'id' column as it has not predictive value

X = df_train.drop(columns=["id", "rainfall"])
y = df_train["rainfall"]

print(X.columns)
print(X.shape)
```

```
Index(['day', 'pressure', 'maxtemp', 'temparature', 'mintemp', 'dewpoint',
       'humidity', 'cloud', 'sunshine', 'winddirection', 'windspeed', 'month'],
      dtype='object')
(2190, 12)
```

```
In [23]: X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)

# Performing a sanity check

print("Train distribution:")
print(y_train.value_counts(normalize=True))

print("\nValidation distribution:")
print(y_val.value_counts(normalize=True))
```

```
Train distribution:
rainfall
1    0.753425
0    0.246575
Name: proportion, dtype: float64

Validation distribution:
rainfall
1    0.753425
0    0.246575
Name: proportion, dtype: float64
```

```
In [24]: for df in [X_train, X_val]:
    if "temparature" in df.columns:
        df.rename(columns={"temparature": "temperature"}, inplace=True)
```

## 7. Column Transformation/Preprocessing Pipeline

```
In [25]: # Step 1: Feature Engineering

class FeatureEngineer(BaseEstimator, TransformerMixin):
    def fit(self, X, y=None):
```

```

# Learn the median for winddirection to avoid leakage
if "winddirection" in X.columns:
    self.winddirection_median_ = X["winddirection"].median()
return self

def transform(self, X):
    X = X.copy()

    # ❶ Fill missing winddirection values with median from training set
    if "winddirection" in X.columns:
        X["winddirection"] = X["winddirection"].fillna(self.winddirection_median_)

    # ❷ Windspeed category using quantiles (Low, Medium, High)
    if "windspeed" in X.columns:
        X["windspeed_category"] = pd.qcut(
            X["windspeed"],
            q=3, # splits data into 3 equal-sized bins
            labels=["Low", "Medium", "High"]
        )

    # ❸ Temperature range
    if {"maxtemp", "mintemp"}.issubset(X.columns):
        X["temp_range"] = X["maxtemp"] - X["mintemp"]

    # ❹ Cyclical day encoding
    if "day" in X.columns:
        X["day_sin"] = np.sin(2 * np.pi * X["day"] / 365)
        X["day_cos"] = np.cos(2 * np.pi * X["day"] / 365)
        X.drop(columns=["day"], inplace=True)

    return X

```

In [26]: # Step 2: Correlation Filter (numeric only)

```

class CorrelationFilter(BaseEstimator, TransformerMixin):
    def __init__(self, threshold=0.9, keep_columns=None):
        self.threshold = threshold
        self.keep_columns = keep_columns
        self.to_drop_ = None

    def fit(self, X, y=None):
        X = pd.DataFrame(X)
        numeric_columns = X.select_dtypes(include=[np.number]).columns
        corr_matrix = X[numeric_columns].corr().abs()
        upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool))
        self.to_drop_ = [col for col in upper.columns if any(upper[col] > self.threshold)]
        if self.keep_columns:
            self.to_drop_ = [c for c in self.to_drop_ if c not in self.keep_columns]
        return self

    def transform(self, X):
        X = pd.DataFrame(X).copy()
        return X.drop(columns=self.to_drop_, errors="ignore")

```

In [27]: # Step 3: Column Transformer (joblib-safe)

```

numeric_features = [
    "windspeed", "temperature", "maxtemp", "mintemp",
    "humidity", "pressure", "dewpoint", "cloud", "sunshine",
]

```

```

        "temp_range", "day_sin", "day_cos"
    ]

categorical_features = [
    "windspeed_category", "winddirection"
]

# Named function for Inf → NaN (joblib-safe)
def inf_to_nan_func(X):
    return np.where(np.isinf(X), np.nan, X)

inf_to_nan = FunctionTransformer(inf_to_nan_func, feature_names_out="one-to-one")

numeric_transformer = Pipeline([
    ("inf_to_nan", inf_to_nan),
    ("imputer", SimpleImputer(strategy="median"))
])

categorical_transformer = Pipeline([
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("encoder", OneHotEncoder(handle_unknown="ignore", sparse_output=False))
])

preprocessor = ColumnTransformer([
    ("num", numeric_transformer, numeric_features),
    ("cat", categorical_transformer, categorical_features)
], remainder="drop")

```

## 8. Pipeline Construction

In [28]: # Step 4: Full Pipeline (preprocessing + MI selection)

```

preprocessing_mi_pipeline = Pipeline([
    ("feature_engineering", FeatureEngineer()),
    ("correlation_filter", CorrelationFilter(threshold=0.9, keep_columns=numeric),
     ("preprocessing", preprocessor), # joblib-safe now
    ("mi_selection", SelectKBest(score_func=mutual_info_classif, k=20))
])

```

## 9. Model Training, Hyperparameter Tuning & Model evaluation

In [29]: # Step 5: Transform

```

X_train_transformed = preprocessing_mi_pipeline.fit_transform(X_train, y_train)
X_val_transformed = preprocessing_mi_pipeline.transform(X_val)

```

In [30]: # Get ALL feature names from ColumnTransformer

```

column_transformer = preprocessing_mi_pipeline.named_steps["preprocessing"]
all_feature_names = column_transformer.get_feature_names_out()

# Apply MI selector mask

mi_selector = preprocessing_mi_pipeline.named_steps["mi_selection"]
selected_mask = mi_selector.get_support()

```

```
selected_feature_names = all_feature_names[selected_mask]

print(f"Total features after MI selection: {len(selected_feature_names)}")
```

Total features after MI selection: 20

In [31]: # 1 Define Optuna objective functions for hyperparameter tuning

```
def rf_objective(trial):
    model = RandomForestClassifier(
        n_estimators=trial.suggest_int("n_estimators", 50, 300),
        max_depth=trial.suggest_int("max_depth", 3, 20),
        min_samples_split=trial.suggest_int("min_samples_split", 2, 20),
        min_samples_leaf=trial.suggest_int("min_samples_leaf", 1, 10),
        random_state=42,
        n_jobs=-1
    )
    cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    return cross_val_score(model, X_train_transformed, y_train, cv=cv, scoring="
```

```
def gb_objective(trial):
    model = GradientBoostingClassifier(
        n_estimators=trial.suggest_int("n_estimators", 50, 300),
        learning_rate=trial.suggest_float("learning_rate", 0.01, 0.3),
        max_depth=trial.suggest_int("max_depth", 2, 10),
        min_samples_split=trial.suggest_int("min_samples_split", 2, 20),
        min_samples_leaf=trial.suggest_int("min_samples_leaf", 1, 10),
        random_state=42
    )
    cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    return cross_val_score(model, X_train_transformed, y_train, cv=cv, scoring="
```

```
def ada_objective(trial):
    model = AdaBoostClassifier(
        n_estimators=trial.suggest_int("n_estimators", 50, 300),
        learning_rate=trial.suggest_float("learning_rate", 0.01, 2.0),
        random_state=42
    )
    cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    return cross_val_score(model, X_train_transformed, y_train, cv=cv, scoring="
```

```
def et_objective(trial):
    model = ExtraTreesClassifier(
        n_estimators=trial.suggest_int("n_estimators", 50, 300),
        max_depth=trial.suggest_int("max_depth", 3, 20),
        min_samples_split=trial.suggest_int("min_samples_split", 2, 20),
        min_samples_leaf=trial.suggest_int("min_samples_leaf", 1, 10),
        random_state=42,
        n_jobs=-1
    )
    cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    return cross_val_score(model, X_train_transformed, y_train, cv=cv, scoring="
```

```
def lr_objective(trial):
    C = trial.suggest_loguniform("C", 1e-3, 1e3)
    penalty = trial.suggest_categorical("penalty", ["l1", "l2"])
    solver = "saga" #  saga supports both L1 and L2
    model = LogisticRegression(
        C=C, penalty=penalty, solver=solver, max_iter=2000, random_state=42
```

```

        )
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
return cross_val_score(model, X_train_transformed, y_train, cv=cv, scoring=""

# 2. Tune models with Optuna

def tune_model(objective_fn, n_trials=30):
    study = optuna.create_study(direction="maximize")
    study.optimize(objective_fn, n_trials=n_trials)
    return study.best_trial.params

best_params = {}
best_params["RandomForest"] = tune_model(rf_objective)
best_params["GradientBoosting"] = tune_model(gb_objective)
best_params["AdaBoost"] = tune_model(ada_objective)
best_params["ExtraTrees"] = tune_model(et_objective)
best_params["LogisticRegression"] = tune_model(lr_objective)

print("\nBest hyperparameters for each model:")
for name, params in best_params.items():
    print(f"{name}: {params}")

# 3. Train final tuned models and evaluate

models = {
    "RandomForest": RandomForestClassifier(**best_params["RandomForest"], random_
    "GradientBoosting": GradientBoostingClassifier(**best_params["GradientBoosti
    "AdaBoost": AdaBoostClassifier(**best_params["AdaBoost"], random_state=42),
    "ExtraTrees": ExtraTreesClassifier(**best_params["ExtraTrees"], random_state
    "LogisticRegression": LogisticRegression(**best_params["LogisticRegression"]
}

for name, model in models.items():
    print(f"\n==== {name} ====")
    model.fit(X_train_transformed, y_train)
    y_pred = model.predict(X_val_transformed)
    y_pred_proba = model.predict_proba(X_val_transformed)[:,1] if hasattr(model,
        accuracy = accuracy_score(y_val, y_pred)
        roc_auc = roc_auc_score(y_val, y_pred_proba) if y_pred_proba is not None else
        report = classification_report(y_val, y_pred)

    print(f"Accuracy: {accuracy:.4f}")
    if roc_auc is not None:
        print(f"ROC AUC: {roc_auc:.4f}")
    print("Classification Report:\n", report)

    # Feature importance
    if hasattr(model, "feature_importances_"):
        fi = pd.DataFrame({
            "Feature": selected_feature_names,
            "Importance": model.feature_importances_
        }).sort_values(by="Importance", ascending=False)
        plt.figure(figsize=(12,6))
        sns.barplot(x="Importance", y="Feature", data=fi, palette="viridis")
        plt.title(f"{name} Feature Importance")
        plt.show()
    elif hasattr(model, "coef_"):
        coefs = abs(model.coef_).flatten()

```

```
fi = pd.DataFrame({
    "Feature": selected_feature_names,
    "Importance": coefs
}).sort_values(by="Importance", ascending=False)
plt.figure(figsize=(12,6))
sns.barplot(x="Importance", y="Feature", data=fi, palette="coolwarm")
plt.title(f"{name} Coefficient Magnitude (Feature Importance)")
plt.show()
```

```
[I 2026-02-14 09:26:59,748] A new study created in memory with name: no-name-295e  
598a-a4ac-4a4f-b991-300e4bc96325  
[I 2026-02-14 09:27:01,379] Trial 0 finished with value: 0.8879735607883162 and p  
arameters: {'n_estimators': 139, 'max_depth': 4, 'min_samples_split': 6, 'min_sa  
mple_leaf': 3}. Best is trial 0 with value: 0.8879735607883162.  
[I 2026-02-14 09:27:02,621] Trial 1 finished with value: 0.8881314143844217 and p  
arameters: {'n_estimators': 91, 'max_depth': 8, 'min_samples_split': 14, 'min_sa  
mple_leaf': 5}. Best is trial 1 with value: 0.8881314143844217.  
[I 2026-02-14 09:27:03,746] Trial 2 finished with value: 0.8872800782484023 and p  
arameters: {'n_estimators': 94, 'max_depth': 4, 'min_samples_split': 10, 'min_sa  
mple_leaf': 10}. Best is trial 1 with value: 0.8881314143844217.  
[I 2026-02-14 09:27:05,193] Trial 3 finished with value: 0.8897592606093008 and p  
arameters: {'n_estimators': 110, 'max_depth': 10, 'min_samples_split': 11, 'min_sa  
mple_leaf': 8}. Best is trial 3 with value: 0.8897592606093008.  
[I 2026-02-14 09:27:07,266] Trial 4 finished with value: 0.8870523600074522 and p  
arameters: {'n_estimators': 153, 'max_depth': 4, 'min_samples_split': 3, 'min_sa  
mple_leaf': 4}. Best is trial 3 with value: 0.8897592606093008.  
[I 2026-02-14 09:27:10,503] Trial 5 finished with value: 0.8884268709549383 and p  
arameters: {'n_estimators': 237, 'max_depth': 4, 'min_samples_split': 14, 'min_sa  
mple_leaf': 4}. Best is trial 3 with value: 0.8897592606093008.  
[I 2026-02-14 09:27:12,895] Trial 6 finished with value: 0.882870302868298 and pa  
rameters: {'n_estimators': 214, 'max_depth': 3, 'min_samples_split': 5, 'min_sa  
mple_leaf': 1}. Best is trial 3 with value: 0.8897592606093008.  
[I 2026-02-14 09:27:15,839] Trial 7 finished with value: 0.8878660299871205 and p  
arameters: {'n_estimators': 234, 'max_depth': 4, 'min_samples_split': 4, 'min_sa  
mple_leaf': 10}. Best is trial 3 with value: 0.8897592606093008.  
[I 2026-02-14 09:27:19,527] Trial 8 finished with value: 0.8897064064866791 and p  
arameters: {'n_estimators': 289, 'max_depth': 19, 'min_samples_split': 14, 'min_sa  
mple_leaf': 10}. Best is trial 3 with value: 0.8897592606093008.  
[I 2026-02-14 09:27:21,215] Trial 9 finished with value: 0.8795258114424114 and p  
arameters: {'n_estimators': 81, 'max_depth': 18, 'min_samples_split': 20, 'min_sa  
mple_leaf': 1}. Best is trial 3 with value: 0.8897592606093008.  
[I 2026-02-14 09:27:22,255] Trial 10 finished with value: 0.8861324755169984 and  
parameters: {'n_estimators': 52, 'max_depth': 13, 'min_samples_split': 9, 'min_sa  
mple_leaf': 7}. Best is trial 3 with value: 0.8897592606093008.  
[I 2026-02-14 09:27:25,924] Trial 11 finished with value: 0.8899612605606991 and  
parameters: {'n_estimators': 272, 'max_depth': 20, 'min_samples_split': 15, 'min_sa  
mple_leaf': 8}. Best is trial 11 with value: 0.8899612605606991.  
[I 2026-02-14 09:27:29,731] Trial 12 finished with value: 0.8891986221477 and pa  
rameters: {'n_estimators': 298, 'max_depth': 13, 'min_samples_split': 19, 'min_sa  
mple_leaf': 7}. Best is trial 11 with value: 0.8899612605606991.  
[I 2026-02-14 09:27:32,190] Trial 13 finished with value: 0.8899017237329186 and  
parameters: {'n_estimators': 187, 'max_depth': 10, 'min_samples_split': 17, 'min_sa  
mple_leaf': 7}. Best is trial 11 with value: 0.8899612605606991.  
[I 2026-02-14 09:27:35,597] Trial 14 finished with value: 0.8893358201096774 and  
parameters: {'n_estimators': 205, 'max_depth': 16, 'min_samples_split': 17, 'min_sa  
mple_leaf': 8}. Best is trial 11 with value: 0.8899612605606991.  
[I 2026-02-14 09:27:39,129] Trial 15 finished with value: 0.8901254931026383 and  
parameters: {'n_estimators': 263, 'max_depth': 8, 'min_samples_split': 17, 'min_sa  
mple_leaf': 6}. Best is trial 15 with value: 0.8901254931026383.  
[I 2026-02-14 09:27:42,355] Trial 16 finished with value: 0.890617684462913 and p  
arameters: {'n_estimators': 261, 'max_depth': 7, 'min_samples_split': 17, 'min_sa  
mple_leaf': 6}. Best is trial 16 with value: 0.890617684462913.  
[I 2026-02-14 09:27:45,611] Trial 17 finished with value: 0.8906175832098047 and  
parameters: {'n_estimators': 262, 'max_depth': 7, 'min_samples_split': 17, 'min_sa  
mple_leaf': 6}. Best is trial 16 with value: 0.890617684462913.  
[I 2026-02-14 09:27:49,949] Trial 18 finished with value: 0.8898538310126121 and  
parameters: {'n_estimators': 260, 'max_depth': 8, 'min_samples_split': 12, 'min_sa  
mple_leaf': 5}. Best is trial 16 with value: 0.890617684462913.  
[I 2026-02-14 09:27:51,937] Trial 19 finished with value: 0.8902533757786364 and
```

```
parameters: {'n_estimators': 158, 'max_depth': 6, 'min_samples_split': 8, 'min_samples_leaf': 6}. Best is trial 16 with value: 0.890617684462913.
[I 2026-02-14 09:27:54,808] Trial 20 finished with value: 0.8902929657440485 and parameters: {'n_estimators': 237, 'max_depth': 6, 'min_samples_split': 19, 'min_samples_leaf': 3}. Best is trial 16 with value: 0.890617684462913.
[I 2026-02-14 09:27:57,718] Trial 21 finished with value: 0.8902929657440485 and parameters: {'n_estimators': 237, 'max_depth': 6, 'min_samples_split': 19, 'min_samples_leaf': 3}. Best is trial 16 with value: 0.890617684462913.
[I 2026-02-14 09:28:01,105] Trial 22 finished with value: 0.8901533377074677 and parameters: {'n_estimators': 249, 'max_depth': 6, 'min_samples_split': 18, 'min_samples_leaf': 2}. Best is trial 16 with value: 0.890617684462913.
[I 2026-02-14 09:28:04,547] Trial 23 finished with value: 0.888910253294776 and parameters: {'n_estimators': 216, 'max_depth': 10, 'min_samples_split': 16, 'min_samples_leaf': 5}. Best is trial 16 with value: 0.890617684462913.
[I 2026-02-14 09:28:07,981] Trial 24 finished with value: 0.8911411630337052 and parameters: {'n_estimators': 281, 'max_depth': 7, 'min_samples_split': 20, 'min_samples_leaf': 4}. Best is trial 24 with value: 0.8911411630337052.
[I 2026-02-14 09:28:11,590] Trial 25 finished with value: 0.8874239589155387 and parameters: {'n_estimators': 283, 'max_depth': 12, 'min_samples_split': 20, 'min_samples_leaf': 4}. Best is trial 24 with value: 0.8911411630337052.
[I 2026-02-14 09:28:16,221] Trial 26 finished with value: 0.8903644504386283 and parameters: {'n_estimators': 298, 'max_depth': 8, 'min_samples_split': 12, 'min_samples_leaf': 6}. Best is trial 24 with value: 0.8911411630337052.
[I 2026-02-14 09:28:18,864] Trial 27 finished with value: 0.8896101147805238 and parameters: {'n_estimators': 190, 'max_depth': 15, 'min_samples_split': 16, 'min_samples_leaf': 9}. Best is trial 24 with value: 0.8911411630337052.
[I 2026-02-14 09:28:22,363] Trial 28 finished with value: 0.8909177986764194 and parameters: {'n_estimators': 274, 'max_depth': 9, 'min_samples_split': 18, 'min_samples_leaf': 5}. Best is trial 24 with value: 0.8911411630337052.
[I 2026-02-14 09:28:26,548] Trial 29 finished with value: 0.8860826589876309 and parameters: {'n_estimators': 275, 'max_depth': 11, 'min_samples_split': 7, 'min_samples_leaf': 4}. Best is trial 24 with value: 0.8911411630337052.
[I 2026-02-14 09:28:26,550] A new study created in memory with name: no-name-71be754c-76a0-478b-9600-f5e94ad240d3
[I 2026-02-14 09:28:34,370] Trial 0 finished with value: 0.8715665070917679 and parameters: {'n_estimators': 144, 'learning_rate': 0.05312478767797437, 'max_depth': 8, 'min_samples_split': 12, 'min_samples_leaf': 5}. Best is trial 0 with value: 0.8715665070917679.
[I 2026-02-14 09:28:38,327] Trial 1 finished with value: 0.8635386543866896 and parameters: {'n_estimators': 149, 'learning_rate': 0.26967300360907565, 'max_depth': 5, 'min_samples_split': 19, 'min_samples_leaf': 6}. Best is trial 0 with value: 0.8715665070917679.
[I 2026-02-14 09:28:45,759] Trial 2 finished with value: 0.8668911448081456 and parameters: {'n_estimators': 300, 'learning_rate': 0.15251212966003974, 'max_depth': 4, 'min_samples_split': 19, 'min_samples_leaf': 5}. Best is trial 0 with value: 0.8715665070917679.
[I 2026-02-14 09:28:47,954] Trial 3 finished with value: 0.8668825382939257 and parameters: {'n_estimators': 67, 'learning_rate': 0.23152748860001304, 'max_depth': 6, 'min_samples_split': 15, 'min_samples_leaf': 7}. Best is trial 0 with value: 0.8715665070917679.
[I 2026-02-14 09:28:52,907] Trial 4 finished with value: 0.8613359942650239 and parameters: {'n_estimators': 293, 'learning_rate': 0.2616221377779615, 'max_depth': 3, 'min_samples_split': 4, 'min_samples_leaf': 6}. Best is trial 0 with value: 0.8715665070917679.
[I 2026-02-14 09:28:58,250] Trial 5 finished with value: 0.8675047386454764 and parameters: {'n_estimators': 96, 'learning_rate': 0.23705237102265775, 'max_depth': 9, 'min_samples_split': 8, 'min_samples_leaf': 10}. Best is trial 0 with value: 0.8715665070917679.
[I 2026-02-14 09:29:06,366] Trial 6 finished with value: 0.8683727815443933 and parameters: {'n_estimators': 256, 'learning_rate': 0.2859695643261104, 'max_depth': 10}.
```

```
h': 6, 'min_samples_split': 20, 'min_samples_leaf': 2}. Best is trial 0 with value: 0.8715665070917679.  
[I 2026-02-14 09:29:14,107] Trial 7 finished with value: 0.8649851562942981 and parameters: {'n_estimators': 248, 'learning_rate': 0.21197884486681876, 'max_depth': 5, 'min_samples_split': 9, 'min_samples_leaf': 5}. Best is trial 0 with value: 0.8715665070917679.  
[I 2026-02-14 09:29:16,275] Trial 8 finished with value: 0.8904951682016637 and parameters: {'n_estimators': 81, 'learning_rate': 0.08794826508985659, 'max_depth': 2, 'min_samples_split': 8, 'min_samples_leaf': 4}. Best is trial 8 with value: 0.8904951682016637.  
[I 2026-02-14 09:29:21,269] Trial 9 finished with value: 0.8692247251990637 and parameters: {'n_estimators': 121, 'learning_rate': 0.14937181663588361, 'max_depth': 7, 'min_samples_split': 17, 'min_samples_leaf': 1}. Best is trial 8 with value: 0.8904951682016637.  
[I 2026-02-14 09:29:24,452] Trial 10 finished with value: 0.8907136217831887 and parameters: {'n_estimators': 192, 'learning_rate': 0.025413994003582185, 'max_depth': 2, 'min_samples_split': 4, 'min_samples_leaf': 2}. Best is trial 10 with value: 0.8907136217831887.  
[I 2026-02-14 09:29:26,911] Trial 11 finished with value: 0.8908965355236406 and parameters: {'n_estimators': 197, 'learning_rate': 0.015321479841586317, 'max_depth': 2, 'min_samples_split': 2, 'min_samples_leaf': 2}. Best is trial 11 with value: 0.8908965355236406.  
[I 2026-02-14 09:29:29,453] Trial 12 finished with value: 0.8909991555490754 and parameters: {'n_estimators': 202, 'learning_rate': 0.014340685898447041, 'max_depth': 2, 'min_samples_split': 2, 'min_samples_leaf': 3}. Best is trial 12 with value: 0.8909991555490754.  
[I 2026-02-14 09:29:32,808] Trial 13 finished with value: 0.8887391861680154 and parameters: {'n_estimators': 190, 'learning_rate': 0.012506087306849276, 'max_depth': 3, 'min_samples_split': 2, 'min_samples_leaf': 3}. Best is trial 12 with value: 0.8909991555490754.  
[I 2026-02-14 09:29:54,759] Trial 14 finished with value: 0.8666684892226192 and parameters: {'n_estimators': 221, 'learning_rate': 0.09485500506482177, 'max_depth': 10, 'min_samples_split': 2, 'min_samples_leaf': 1}. Best is trial 12 with value: 0.8909991555490754.  
[I 2026-02-14 09:29:58,482] Trial 15 finished with value: 0.8802023847132107 and parameters: {'n_estimators': 218, 'learning_rate': 0.0724242077121705, 'max_depth': 3, 'min_samples_split': 5, 'min_samples_leaf': 3}. Best is trial 12 with value: 0.8909991555490754.  
[I 2026-02-14 09:30:00,522] Trial 16 finished with value: 0.8874536260763207 and parameters: {'n_estimators': 166, 'learning_rate': 0.1232268652795309, 'max_depth': 2, 'min_samples_split': 6, 'min_samples_leaf': 8}. Best is trial 12 with value: 0.8909991555490754.  
[I 2026-02-14 09:30:06,663] Trial 17 finished with value: 0.8806144848646854 and parameters: {'n_estimators': 233, 'learning_rate': 0.0425827836677503, 'max_depth': 4, 'min_samples_split': 11, 'min_samples_leaf': 3}. Best is trial 12 with value: 0.8909991555490754.  
[I 2026-02-14 09:30:10,885] Trial 18 finished with value: 0.8699286368091501 and parameters: {'n_estimators': 188, 'learning_rate': 0.19618061503036865, 'max_depth': 4, 'min_samples_split': 2, 'min_samples_leaf': 2}. Best is trial 12 with value: 0.8909991555490754.  
[I 2026-02-14 09:30:14,293] Trial 19 finished with value: 0.8914468967947317 and parameters: {'n_estimators': 273, 'learning_rate': 0.011782613983568804, 'max_depth': 2, 'min_samples_split': 13, 'min_samples_leaf': 4}. Best is trial 19 with value: 0.8914468967947317.  
[I 2026-02-14 09:30:19,908] Trial 20 finished with value: 0.8810142321369266 and parameters: {'n_estimators': 274, 'learning_rate': 0.05895875067175775, 'max_depth': 3, 'min_samples_split': 13, 'min_samples_leaf': 4}. Best is trial 19 with value: 0.8914468967947317.  
[I 2026-02-14 09:30:22,483] Trial 21 finished with value: 0.8913836136019377 and parameters: {'n_estimators': 207, 'learning_rate': 0.02177478458702009, 'max_depth': 2}
```

```
h': 2, 'min_samples_split': 15, 'min_samples_leaf': 4}. Best is trial 19 with value: 0.8914468967947317.  
[I 2026-02-14 09:30:25,519] Trial 22 finished with value: 0.8911259244408803 and parameters: {'n_estimators': 245, 'learning_rate': 0.03553978054743935, 'max_depth': 2, 'min_samples_split': 15, 'min_samples_leaf': 4}. Best is trial 19 with value: 0.8914468967947317.  
[I 2026-02-14 09:30:30,657] Trial 23 finished with value: 0.8738510809781861 and parameters: {'n_estimators': 266, 'learning_rate': 0.10969219028240622, 'max_depth': 3, 'min_samples_split': 15, 'min_samples_leaf': 4}. Best is trial 19 with value: 0.8914468967947317.  
[I 2026-02-14 09:30:37,582] Trial 24 finished with value: 0.8784354167172932 and parameters: {'n_estimators': 239, 'learning_rate': 0.03979217607972785, 'max_depth': 5, 'min_samples_split': 14, 'min_samples_leaf': 4}. Best is trial 19 with value: 0.8914468967947317.  
[I 2026-02-14 09:30:44,649] Trial 25 finished with value: 0.8735815452034377 and parameters: {'n_estimators': 284, 'learning_rate': 0.06756170082157004, 'max_depth': 4, 'min_samples_split': 17, 'min_samples_leaf': 7}. Best is trial 19 with value: 0.8914468967947317.  
[I 2026-02-14 09:30:48,117] Trial 26 finished with value: 0.8902659817906411 and parameters: {'n_estimators': 264, 'learning_rate': 0.03594473448759752, 'max_depth': 2, 'min_samples_split': 17, 'min_samples_leaf': 6}. Best is trial 19 with value: 0.8914468967947317.  
[I 2026-02-14 09:30:51,759] Trial 27 finished with value: 0.8715930354061869 and parameters: {'n_estimators': 216, 'learning_rate': 0.1291803576633805, 'max_depth': 3, 'min_samples_split': 15, 'min_samples_leaf': 5}. Best is trial 19 with value: 0.8914468967947317.  
[I 2026-02-14 09:31:01,920] Trial 28 finished with value: 0.873231816966781 and parameters: {'n_estimators': 230, 'learning_rate': 0.09669604397864989, 'max_depth': 7, 'min_samples_split': 10, 'min_samples_leaf': 4}. Best is trial 19 with value: 0.8914468967947317.  
[I 2026-02-14 09:31:09,590] Trial 29 finished with value: 0.8726874802556438 and parameters: {'n_estimators': 162, 'learning_rate': 0.05256322592116986, 'max_depth': 8, 'min_samples_split': 12, 'min_samples_leaf': 5}. Best is trial 19 with value: 0.8914468967947317.  
[I 2026-02-14 09:31:09,593] A new study created in memory with name: no-name-475b  
cf33-2b93-49df-a4bc-de007a38f61f  
[I 2026-02-14 09:31:11,326] Trial 0 finished with value: 0.8898573748714085 and parameters: {'n_estimators': 93, 'learning_rate': 0.08984868829153067}. Best is trial 0 with value: 0.8898573748714085.  
[I 2026-02-14 09:31:15,326] Trial 1 finished with value: 0.8764529314799964 and parameters: {'n_estimators': 283, 'learning_rate': 1.573089806165302}. Best is trial 0 with value: 0.8898573748714085.  
[I 2026-02-14 09:31:18,703] Trial 2 finished with value: 0.8880942544936129 and parameters: {'n_estimators': 258, 'learning_rate': 0.4498666503745843}. Best is trial 0 with value: 0.8898573748714085.  
[I 2026-02-14 09:31:21,878] Trial 3 finished with value: 0.8758105817598599 and parameters: {'n_estimators': 246, 'learning_rate': 1.6037865913786884}. Best is trial 0 with value: 0.8898573748714085.  
[I 2026-02-14 09:31:24,943] Trial 4 finished with value: 0.8589140705369654 and parameters: {'n_estimators': 172, 'learning_rate': 1.9070672641097264}. Best is trial 0 with value: 0.8898573748714085.  
[I 2026-02-14 09:31:27,349] Trial 5 finished with value: 0.8910924096619766 and parameters: {'n_estimators': 161, 'learning_rate': 0.13811002440205775}. Best is trial 5 with value: 0.8910924096619766.  
[I 2026-02-14 09:31:28,937] Trial 6 finished with value: 0.8809393048366585 and parameters: {'n_estimators': 123, 'learning_rate': 1.3694310436648687}. Best is trial 5 with value: 0.8910924096619766.  
[I 2026-02-14 09:31:29,713] Trial 7 finished with value: 0.8814752375397925 and parameters: {'n_estimators': 58, 'learning_rate': 1.7615904145646333}. Best is trial 5 with value: 0.8910924096619766.
```

```
[I 2026-02-14 09:31:31,912] Trial 8 finished with value: 0.8865093902132795 and parameters: {'n_estimators': 169, 'learning_rate': 0.551544886053529}. Best is trial 5 with value: 0.8910924096619766.  
[I 2026-02-14 09:31:34,111] Trial 9 finished with value: 0.8583985403351884 and parameters: {'n_estimators': 168, 'learning_rate': 1.9818214775841065}. Best is trial 5 with value: 0.8910924096619766.  
[I 2026-02-14 09:31:37,375] Trial 10 finished with value: 0.8863179205851619 and parameters: {'n_estimators': 215, 'learning_rate': 1.0081703346947195}. Best is trial 5 with value: 0.8910924096619766.  
[I 2026-02-14 09:31:39,202] Trial 11 finished with value: 0.8837680635545512 and parameters: {'n_estimators': 96, 'learning_rate': 0.022554986559257695}. Best is trial 5 with value: 0.8910924096619766.  
[I 2026-02-14 09:31:40,873] Trial 12 finished with value: 0.8895430852227163 and parameters: {'n_estimators': 125, 'learning_rate': 0.05177329978943576}. Best is trial 5 with value: 0.8910924096619766.  
[I 2026-02-14 09:31:41,558] Trial 13 finished with value: 0.8886813200165244 and parameters: {'n_estimators': 51, 'learning_rate': 0.3809071875808233}. Best is trial 5 with value: 0.8910924096619766.  
[I 2026-02-14 09:31:42,912] Trial 14 finished with value: 0.8818303828177525 and parameters: {'n_estimators': 105, 'learning_rate': 0.7576107888574175}. Best is trial 5 with value: 0.8910924096619766.  
[I 2026-02-14 09:31:44,716] Trial 15 finished with value: 0.8912070281807651 and parameters: {'n_estimators': 138, 'learning_rate': 0.18265124489785745}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:31:47,310] Trial 16 finished with value: 0.8859683948547221 and parameters: {'n_estimators': 197, 'learning_rate': 0.98061711149807}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:31:49,175] Trial 17 finished with value: 0.8908089515848137 and parameters: {'n_estimators': 143, 'learning_rate': 0.29073120053174706}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:31:52,951] Trial 18 finished with value: 0.8880462605201981 and parameters: {'n_estimators': 210, 'learning_rate': 0.6987899833178266}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:31:54,818] Trial 19 finished with value: 0.891005129482475 and parameters: {'n_estimators': 144, 'learning_rate': 0.22436922812618937}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:31:55,871] Trial 20 finished with value: 0.8838102861007833 and parameters: {'n_estimators': 80, 'learning_rate': 1.042977635551596}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:31:57,758] Trial 21 finished with value: 0.8904864604343354 and parameters: {'n_estimators': 144, 'learning_rate': 0.24065961458597362}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:31:59,759] Trial 22 finished with value: 0.8909838157031421 and parameters: {'n_estimators': 145, 'learning_rate': 0.21278484260653185}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:32:02,223] Trial 23 finished with value: 0.8876843819105247 and parameters: {'n_estimators': 190, 'learning_rate': 0.5481607493792615}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:32:04,296] Trial 24 finished with value: 0.8828225114011001 and parameters: {'n_estimators': 123, 'learning_rate': 0.7488866898678586}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:32:06,926] Trial 25 finished with value: 0.8906851696597086 and parameters: {'n_estimators': 156, 'learning_rate': 0.1761562776891065}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:32:09,811] Trial 26 finished with value: 0.8906687666561363 and parameters: {'n_estimators': 222, 'learning_rate': 0.36762996118706637}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:32:12,252] Trial 27 finished with value: 0.8874993924813491 and parameters: {'n_estimators': 189, 'learning_rate': 0.5483219882197348}. Best is trial 15 with value: 0.8912070281807651.
```

```
[I 2026-02-14 09:32:13,771] Trial 28 finished with value: 0.8905594132989882 and parameters: {'n_estimators': 117, 'learning_rate': 0.14151962884447256}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:32:14,911] Trial 29 finished with value: 0.8637972548257231 and parameters: {'n_estimators': 87, 'learning_rate': 0.012485967379583962}. Best is trial 15 with value: 0.8912070281807651.  
[I 2026-02-14 09:32:14,912] A new study created in memory with name: no-name-5eae3cbd-b404-4689-9b33-576b04555a06  
[I 2026-02-14 09:32:15,658] Trial 0 finished with value: 0.8930459365102509 and parameters: {'n_estimators': 82, 'max_depth': 10, 'min_samples_split': 7, 'min_samples_leaf': 7}. Best is trial 0 with value: 0.8930459365102509.  
[I 2026-02-14 09:32:18,515] Trial 1 finished with value: 0.8910137866232493 and parameters: {'n_estimators': 285, 'max_depth': 17, 'min_samples_split': 12, 'min_samples_leaf': 10}. Best is trial 0 with value: 0.8930459365102509.  
[I 2026-02-14 09:32:20,608] Trial 2 finished with value: 0.8921106615473094 and parameters: {'n_estimators': 219, 'max_depth': 6, 'min_samples_split': 11, 'min_samples_leaf': 4}. Best is trial 0 with value: 0.8930459365102509.  
[I 2026-02-14 09:32:21,322] Trial 3 finished with value: 0.8916907649064827 and parameters: {'n_estimators': 80, 'max_depth': 17, 'min_samples_split': 18, 'min_samples_leaf': 8}. Best is trial 0 with value: 0.8930459365102509.  
[I 2026-02-14 09:32:22,232] Trial 4 finished with value: 0.8926309000186305 and parameters: {'n_estimators': 97, 'max_depth': 13, 'min_samples_split': 16, 'min_samples_leaf': 6}. Best is trial 0 with value: 0.8930459365102509.  
[I 2026-02-14 09:32:24,471] Trial 5 finished with value: 0.8912430236608264 and parameters: {'n_estimators': 273, 'max_depth': 10, 'min_samples_split': 10, 'min_samples_leaf': 10}. Best is trial 0 with value: 0.8930459365102509.  
[I 2026-02-14 09:32:25,087] Trial 6 finished with value: 0.887565409508072 and parameters: {'n_estimators': 58, 'max_depth': 16, 'min_samples_split': 19, 'min_samples_leaf': 1}. Best is trial 0 with value: 0.8930459365102509.  
[I 2026-02-14 09:32:25,603] Trial 7 finished with value: 0.8917124330716953 and parameters: {'n_estimators': 52, 'max_depth': 13, 'min_samples_split': 6, 'min_samples_leaf': 9}. Best is trial 0 with value: 0.8930459365102509.  
[I 2026-02-14 09:32:26,326] Trial 8 finished with value: 0.8905940924886394 and parameters: {'n_estimators': 75, 'max_depth': 13, 'min_samples_split': 6, 'min_samples_leaf': 3}. Best is trial 0 with value: 0.8930459365102509.  
[I 2026-02-14 09:32:26,912] Trial 9 finished with value: 0.8852563323694037 and parameters: {'n_estimators': 66, 'max_depth': 3, 'min_samples_split': 9, 'min_samples_leaf': 7}. Best is trial 0 with value: 0.8930459365102509.  
[I 2026-02-14 09:32:28,082] Trial 10 finished with value: 0.8922186986140475 and parameters: {'n_estimators': 137, 'max_depth': 8, 'min_samples_split': 2, 'min_samples_leaf': 5}. Best is trial 0 with value: 0.8930459365102509.  
[I 2026-02-14 09:32:29,206] Trial 11 finished with value: 0.8921519728155655 and parameters: {'n_estimators': 126, 'max_depth': 20, 'min_samples_split': 15, 'min_samples_leaf': 6}. Best is trial 0 with value: 0.8930459365102509.  
[I 2026-02-14 09:32:30,497] Trial 12 finished with value: 0.8931797931196488 and parameters: {'n_estimators': 130, 'max_depth': 10, 'min_samples_split': 16, 'min_samples_leaf': 7}. Best is trial 12 with value: 0.8931797931196488.  
[I 2026-02-14 09:32:32,438] Trial 13 finished with value: 0.8919987768624497 and parameters: {'n_estimators': 174, 'max_depth': 9, 'min_samples_split': 14, 'min_samples_leaf': 8}. Best is trial 12 with value: 0.8931797931196488.  
[I 2026-02-14 09:32:33,679] Trial 14 finished with value: 0.8907570087401684 and parameters: {'n_estimators': 125, 'max_depth': 6, 'min_samples_split': 7, 'min_samples_leaf': 7}. Best is trial 12 with value: 0.8931797931196488.  
[I 2026-02-14 09:32:35,155] Trial 15 finished with value: 0.8929799194835282 and parameters: {'n_estimators': 169, 'max_depth': 10, 'min_samples_split': 2, 'min_samples_leaf': 4}. Best is trial 12 with value: 0.8931797931196488.  
[I 2026-02-14 09:32:36,897] Trial 16 finished with value: 0.8916504661693114 and parameters: {'n_estimators': 211, 'max_depth': 7, 'min_samples_split': 13, 'min_samples_leaf': 8}. Best is trial 12 with value: 0.8931797931196488.  
[I 2026-02-14 09:32:37,774] Trial 17 finished with value: 0.8866522077227772 and
```

```
parameters: {'n_estimators': 102, 'max_depth': 4, 'min_samples_split': 17, 'min_samples_leaf': 7}. Best is trial 12 with value: 0.8931797931196488.  
[I 2026-02-14 09:32:39,089] Trial 18 finished with value: 0.8908772974330311 and parameters: {'n_estimators': 149, 'max_depth': 11, 'min_samples_split': 20, 'min_samples_leaf': 5}. Best is trial 12 with value: 0.8931797931196488.  
[I 2026-02-14 09:32:41,143] Trial 19 finished with value: 0.8872149724996558 and parameters: {'n_estimators': 216, 'max_depth': 15, 'min_samples_split': 8, 'min_samples_leaf': 1}. Best is trial 12 with value: 0.8931797931196488.  
[I 2026-02-14 09:32:41,987] Trial 20 finished with value: 0.8883762444007031 and parameters: {'n_estimators': 100, 'max_depth': 5, 'min_samples_split': 4, 'min_samples_leaf': 9}. Best is trial 12 with value: 0.8931797931196488.  
[I 2026-02-14 09:32:43,577] Trial 21 finished with value: 0.8927349882141382 and parameters: {'n_estimators': 170, 'max_depth': 10, 'min_samples_split': 2, 'min_samples_leaf': 4}. Best is trial 12 with value: 0.8931797931196488.  
[I 2026-02-14 09:32:45,858] Trial 22 finished with value: 0.8931918422395567 and parameters: {'n_estimators': 193, 'max_depth': 11, 'min_samples_split': 4, 'min_samples_leaf': 3}. Best is trial 22 with value: 0.8931918422395567.  
[I 2026-02-14 09:32:47,895] Trial 23 finished with value: 0.8917955618737496 and parameters: {'n_estimators': 198, 'max_depth': 12, 'min_samples_split': 4, 'min_samples_leaf': 3}. Best is trial 22 with value: 0.8931918422395567.  
[I 2026-02-14 09:32:50,060] Trial 24 finished with value: 0.8924960308781479 and parameters: {'n_estimators': 252, 'max_depth': 8, 'min_samples_split': 4, 'min_samples_leaf': 2}. Best is trial 22 with value: 0.8931918422395567.  
[I 2026-02-14 09:32:51,378] Trial 25 finished with value: 0.8924331526977879 and parameters: {'n_estimators': 150, 'max_depth': 14, 'min_samples_split': 6, 'min_samples_leaf': 6}. Best is trial 22 with value: 0.8931918422395567.  
[I 2026-02-14 09:32:53,103] Trial 26 finished with value: 0.8919423788810317 and parameters: {'n_estimators': 194, 'max_depth': 12, 'min_samples_split': 8, 'min_samples_leaf': 5}. Best is trial 22 with value: 0.8931918422395567.  
[I 2026-02-14 09:32:55,141] Trial 27 finished with value: 0.8932708196641638 and parameters: {'n_estimators': 238, 'max_depth': 11, 'min_samples_split': 10, 'min_samples_leaf': 7}. Best is trial 27 with value: 0.8932708196641638.  
[I 2026-02-14 09:32:57,370] Trial 28 finished with value: 0.8919425813872486 and parameters: {'n_estimators': 245, 'max_depth': 8, 'min_samples_split': 12, 'min_samples_leaf': 3}. Best is trial 27 with value: 0.8932708196641638.  
[I 2026-02-14 09:33:00,037] Trial 29 finished with value: 0.8916114837225504 and parameters: {'n_estimators': 248, 'max_depth': 11, 'min_samples_split': 15, 'min_samples_leaf': 9}. Best is trial 27 with value: 0.8932708196641638.  
[I 2026-02-14 09:33:00,039] A new study created in memory with name: no-name-2889a58f-4fdf-418b-9d88-d33c7a283fb7  
[I 2026-02-14 09:33:04,928] Trial 0 finished with value: 0.8951055259896478 and parameters: {'C': 0.634900553321891, 'penalty': 'l2'}. Best is trial 0 with value: 0.8951055259896478.  
[I 2026-02-14 09:33:05,870] Trial 1 finished with value: 0.8860337537362397 and parameters: {'C': 0.045768741532036215, 'penalty': 'l1'}. Best is trial 0 with value: 0.8951055259896478.  
[I 2026-02-14 09:33:07,752] Trial 2 finished with value: 0.8920399868775972 and parameters: {'C': 0.03303057646666202, 'penalty': 'l2'}. Best is trial 0 with value: 0.8951055259896478.  
[I 2026-02-14 09:33:15,202] Trial 3 finished with value: 0.8949560764015455 and parameters: {'C': 436.1034259701939, 'penalty': 'l1'}. Best is trial 0 with value: 0.8951055259896478.  
[I 2026-02-14 09:33:20,146] Trial 4 finished with value: 0.8949914137364017 and parameters: {'C': 0.71488295752957, 'penalty': 'l2'}. Best is trial 0 with value: 0.8951055259896478.  
[I 2026-02-14 09:33:20,526] Trial 5 finished with value: 0.886289519088236 and parameters: {'C': 0.003858872298230565, 'penalty': 'l2'}. Best is trial 0 with value: 0.8951055259896478.  
[I 2026-02-14 09:33:20,930] Trial 6 finished with value: 0.8850811644917499 and parameters: {'C': 0.028335668997371267, 'penalty': 'l1'}. Best is trial 0 with val
```

```
ue: 0.8951055259896478.  
[I 2026-02-14 09:33:28,279] Trial 7 finished with value: 0.8951577725936186 and parameters: {'C': 5.7805221265016895, 'penalty': 'l1'}. Best is trial 7 with value: 0.8951577725936186.  
[I 2026-02-14 09:33:28,733] Trial 8 finished with value: 0.8866945315221176 and parameters: {'C': 0.004661070310171968, 'penalty': 'l2'}. Best is trial 7 with value: 0.8951577725936186.  
[I 2026-02-14 09:33:34,426] Trial 9 finished with value: 0.8950876041894487 and parameters: {'C': 3.1472738561516707, 'penalty': 'l1'}. Best is trial 7 with value: 0.8951577725936186.  
[I 2026-02-14 09:33:41,934] Trial 10 finished with value: 0.8949473686342172 and parameters: {'C': 57.76301047698917, 'penalty': 'l1'}. Best is trial 7 with value: 0.8951577725936186.  
[I 2026-02-14 09:33:46,803] Trial 11 finished with value: 0.8950085255117333 and parameters: {'C': 3.119761308196007, 'penalty': 'l2'}. Best is trial 7 with value: 0.8951577725936186.  
[I 2026-02-14 09:33:52,691] Trial 12 finished with value: 0.8949472673811085 and parameters: {'C': 29.146538656976315, 'penalty': 'l2'}. Best is trial 7 with value: 0.8951577725936186.  
[I 2026-02-14 09:33:58,225] Trial 13 finished with value: 0.8937358751913684 and parameters: {'C': 0.3626380316661745, 'penalty': 'l1'}. Best is trial 7 with value: 0.8951577725936186.  
[I 2026-02-14 09:34:04,414] Trial 14 finished with value: 0.8950523681077008 and parameters: {'C': 11.622855049139504, 'penalty': 'l1'}. Best is trial 7 with value: 0.8951577725936186.  
[I 2026-02-14 09:34:10,087] Trial 15 finished with value: 0.8947920463658233 and parameters: {'C': 0.26928769872502517, 'penalty': 'l2'}. Best is trial 7 with value: 0.8951577725936186.  
[I 2026-02-14 09:34:14,930] Trial 16 finished with value: 0.8949560764015455 and parameters: {'C': 394.2151168765312, 'penalty': 'l2'}. Best is trial 7 with value: 0.8951577725936186.  
[I 2026-02-14 09:34:22,297] Trial 17 finished with value: 0.8951137274914341 and parameters: {'C': 7.2536571005000425, 'penalty': 'l1'}. Best is trial 7 with value: 0.8951577725936186.  
[I 2026-02-14 09:34:28,324] Trial 18 finished with value: 0.8949473686342172 and parameters: {'C': 58.27036634429128, 'penalty': 'l1'}. Best is trial 7 with value: 0.8951577725936186.  
[I 2026-02-14 09:34:35,698] Trial 19 finished with value: 0.8951667841202726 and parameters: {'C': 4.308661382633495, 'penalty': 'l1'}. Best is trial 19 with value: 0.8951667841202726.  
[I 2026-02-14 09:34:38,569] Trial 20 finished with value: 0.8906109005046454 and parameters: {'C': 0.1435945661074243, 'penalty': 'l1'}. Best is trial 19 with value: 0.8951667841202726.  
[I 2026-02-14 09:34:44,823] Trial 21 finished with value: 0.8951229415243048 and parameters: {'C': 3.6950872878853023, 'penalty': 'l1'}. Best is trial 19 with value: 0.8951667841202726.  
[I 2026-02-14 09:34:51,778] Trial 22 finished with value: 0.8951052222303225 and parameters: {'C': 3.2742108980554674, 'penalty': 'l1'}. Best is trial 19 with value: 0.8951667841202726.  
[I 2026-02-14 09:34:57,887] Trial 23 finished with value: 0.8950084242586248 and parameters: {'C': 21.141022508016658, 'penalty': 'l1'}. Best is trial 19 with value: 0.8951667841202726.  
[I 2026-02-14 09:35:05,329] Trial 24 finished with value: 0.8949560764015455 and parameters: {'C': 136.91256441087077, 'penalty': 'l1'}. Best is trial 19 with value: 0.8951667841202726.  
[I 2026-02-14 09:35:11,031] Trial 25 finished with value: 0.8950623921654394 and parameters: {'C': 1.7404465400336795, 'penalty': 'l1'}. Best is trial 19 with value: 0.8951667841202726.  
[I 2026-02-14 09:35:18,408] Trial 26 finished with value: 0.8950522668545926 and parameters: {'C': 10.033346241683715, 'penalty': 'l1'}. Best is trial 19 with val
```

```

ue: 0.8951667841202726.
[I 2026-02-14 09:35:23,984] Trial 27 finished with value: 0.8948007541331519 and
parameters: {'C': 1.0442243816127308, 'penalty': 'l1'}. Best is trial 19 with val
ue: 0.8951667841202726.
[I 2026-02-14 09:35:27,516] Trial 28 finished with value: 0.8884338574194228 and
parameters: {'C': 0.1012445562390169, 'penalty': 'l1'}. Best is trial 19 with val
ue: 0.8951667841202726.
[I 2026-02-14 09:35:33,553] Trial 29 finished with value: 0.8949560764015455 and
parameters: {'C': 133.6640919857623, 'penalty': 'l1'}. Best is trial 19 with val
ue: 0.8951667841202726.

```

Best hyperparameters for each model:

```

RandomForest: {'n_estimators': 281, 'max_depth': 7, 'min_samples_split': 20, 'min
_samples_leaf': 4}
GradientBoosting: {'n_estimators': 273, 'learning_rate': 0.011782613983568804, 'm
ax_depth': 2, 'min_samples_split': 13, 'min_samples_leaf': 4}
AdaBoost: {'n_estimators': 138, 'learning_rate': 0.18265124489785745}
ExtraTrees: {'n_estimators': 238, 'max_depth': 11, 'min_samples_split': 10, 'min_
samples_leaf': 7}
LogisticRegression: {'C': 4.308661382633495, 'penalty': 'l1'}

```


===== RandomForest =====

Accuracy: 0.8767

ROC AUC: 0.8796

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.65	0.72	108
1	0.89	0.95	0.92	330
accuracy			0.88	438
macro avg	0.85	0.80	0.82	438
weighted avg	0.87	0.88	0.87	438




===== GradientBoosting =====

Accuracy: 0.8699

ROC AUC: 0.8753

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.62	0.70	108
1	0.88	0.95	0.92	330
accuracy			0.87	438
macro avg	0.85	0.79	0.81	438
weighted avg	0.87	0.87	0.86	438




===== AdaBoost =====

Accuracy: 0.8676

ROC AUC: 0.8774

Classification Report:

	precision	recall	f1-score	support
0	0.77	0.67	0.71	108
1	0.90	0.93	0.91	330
accuracy			0.87	438
macro avg	0.83	0.80	0.81	438
weighted avg	0.86	0.87	0.86	438




===== ExtraTrees =====

Accuracy: 0.8653

ROC AUC: 0.8789

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.59	0.68	108
1	0.88	0.95	0.91	330
accuracy			0.87	438
macro avg	0.84	0.77	0.80	438
weighted avg	0.86	0.87	0.86	438




===== LogisticRegression =====

Accuracy: 0.8630

ROC AUC: 0.8678

Classification Report:

	precision	recall	f1-score	support
0	0.78	0.62	0.69	108
1	0.88	0.94	0.91	330
accuracy			0.86	438
macro avg	0.83	0.78	0.80	438
weighted avg	0.86	0.86	0.86	438



```
In [32]: # STEP 4: Compare all models & select best
```

```
results = []
trained_models = {}

for name, model in models.items():
    # store trained model
    trained_models[name] = model

    # predictions
    y_pred = model.predict(X_val_transformed)
    if hasattr(model, "predict_proba"):
        y_pred_proba = model.predict_proba(X_val_transformed)[:, 1]
        roc_auc = roc_auc_score(y_val, y_pred_proba)
    else:
        roc_auc = None
    accuracy = accuracy_score(y_val, y_pred)

    # append results
    results.append({
        "Model": name,
        "Accuracy": accuracy,
        "ROC_AUC": roc_auc
    })

# Convert to DataFrame
results_df = pd.DataFrame(results).sort_values(by="ROC_AUC", ascending=False).reset_index()
display(results_df)

# Select best model
best_model_name = results_df.loc[0, "Model"]
best_model = trained_models[best_model_name]
print(f"Best model selected: {best_model_name}")
```

	Model	Accuracy	ROC_AUC
0	RandomForest	0.876712	0.879602
1	ExtraTrees	0.865297	0.878872
2	AdaBoost	0.867580	0.877357
3	GradientBoosting	0.869863	0.875309
4	LogisticRegression	0.863014	0.867789

Best model selected: RandomForest

## 10. Model Persistence/ Deployment

```
In [33]: final_pipeline = Pipeline([
    ("preprocessing", preprocessing_mi_pipeline), # your feature engineering +
    ("model", best_model) # the best trained model
])
```

```
In [34]: # Save full pipeline (preprocessing + SMOTE + model)
```

```
MODEL_PATH = f"best_pipeline_{best_model_name}.joblib"
```

```
joblib.dump(final_pipeline, MODEL_PATH)
print(f"Pipeline saved to {MODEL_PATH}")
```

Pipeline saved to best\_pipeline\_RandomForest.joblib

```
In [35]: # Fix column typo
if "temparature" in df_test.columns:
    df_test.rename(columns={"temparature": "temperature"}, inplace=True)
```

```
In [36]: print(df_test.columns)
```

```
Index(['id', 'day', 'pressure', 'maxtemp', 'temperature', 'mintemp',
       'dewpoint', 'humidity', 'cloud', 'sunshine', 'winddirection',
       'windspeed'],
      dtype='object')
```

```
In [37]: X_test = df_test # no need to drop anything
```

```
In [38]: loaded_pipeline = joblib.load(MODEL_PATH)

# Predicted class labels (rainfall or no rainfall)
y_pred = loaded_pipeline.predict(X_test)

# Predicted probabilities (for the positive class, e.g., rainfall)
y_proba = loaded_pipeline.predict_proba(X_test)[:, 1]
```

```
In [39]: # Save predictions to CSV
output_df = X_test.copy()
output_df["predicted_rainfall"] = y_pred
output_df["rainfall_prob"] = y_proba

output_df.to_csv("rainfall_predictions.csv", index=False)
print("Predictions saved to rainfall_predictions.csv")
```

Predictions saved to rainfall\_predictions.csv